

milestone_2_00951537

June 16, 2024

1 Preamble

See the script below for all package imports

```
[1]: %run -i scripts/preamble.py
```

```
2024-06-16 20:50:35.684937: I tensorflow/core/platform/cpu_feature_guard.cc:182]
This TensorFlow binary is optimized to use available CPU instructions in
performance-critical operations.
To enable the following instructions: SSE4.1 SSE4.2 AVX AVX2 AVX512F FMA, in
other operations, rebuild TensorFlow with the appropriate compiler flags.
```

Note: Due to submission size limitations, the datasets were **not** included in the repost submission. Instead, I have made them publicly downloadable from S3. Please refer to the data download instructions in the README file.

2 Introduction

The global digital advertising market is worth approximately \$602 billion today. Due to the increasing rate of online participation since the COVID-19 pandemic, this number has been rapidly increasing and is expected to reach \$871 billion by the end of 2027 (eMarketer, 2024). Many of the major Ad platforms such as Google, Facebook and Amazon operate on a cost-per-user-engagement pricing model, which usually means that advertisers get charged for every time a user clicks on an advertisement. This means that these platforms are incentivized to make sure that the content shown to each user is as relevant as possible in order to maximize the number of clicks in the long term. Attaining accurate Click-Through Rate (CTR) prediction is a necessary first step for Ad personalization, which is why study of CTR prediction methods have been an extremely active part of Machine Learning research over the past through years.

Initially, shallow prediction methods such as XGBoost (Cite), Factorization Machines (Cite) and Field-Aware Factorization Machines (Cite) have been used for CTR prediction. However, these methods have often been shown to be unable to capture the higher order feature interactions in the sparse multy value categorical Ad Marketplace datasets (Cite). Since then, Deep Learning methods have been shown to show superior predictive ability on these datasets. The focus of my reasearch project is therefore to explore the merits of different Deep Learning architechtures for click-through rate prediction

In the following report, I explore the relevant datasets and simulations that I will be using throughout my research project. In the first section, I perform an exploratory data analysis on three widely

adopted benchmark CTR prediction datasets; the KDD12 (Aden, 2012), Avazu (Wang and Cukierski, 2014) and Criteo (Tien et al, 2014) datasets. In the second section, I then explore possible ways of simulating the ad marketplace environment in order to test the reinforcement learning framework.

3 Data Analysis and Pre-processing

I begin below by first introducing the three datasets widely used as benchmarks in CTR prediction research.

3.0.1 KDD12

The **KDD12** dataset was first released for the KDD Cup 2012 competition (Cite), with the original task being to predict the number of clicks for a given number of impressions. Each line represents a training instance derived from the session logs for the advertizing marketplace. In the context of this dataset, a “session” refers to an interaction between a user and the search engine, containing the following components; the user, a list of adverts returned by the search engine and shown (impressed) to the user and zero or more adverts clicked on by the user. Each line in the training set includes:

- **Click and Impression counts:** The click counts were the original target variable when the dataset was first released for the competition. As done in (Cite Song and Others), this dataset can be adapted to CTR prediction by simply calculating the CTR for each instance by dividing the Click counts by the Impression counts.
- **Session features:** These include *session depth* (the number of ads impressed in a session) as well as the tokenized query phrase that the user entered into the search engine.
- **User features:** Encoded gender and age group for the user, if known.
- **Ad features:** Display URL, ad ID, advertiser ID and encoded title, description and purchased key words.

```
[2]: %run -i scripts/data_analysis_and_preprocessing/retrieve_kdd12.py
```

Snapshot of KDD12 training data:

	Click	Impression	DisplayURL	AdID	AdvertiserID	Depth	\
0	0	1	12057878999086460853	20157098	27961	1	
1	0	1	12057878999086460853	20221208	27961	2	
2	0	1	12057878999086460853	20183701	27961	1	
3	0	1	12057878999086460853	20183690	27961	1	
4	0	1	3029113635936639912	10397010	24973	2	

	Position	QueryID	KeywordID	TitleID	DescriptionID	UserID
0	1	75606	15055	12391	13532	1350148
1	1	2977	1278	3054	4561	1350148
2	1	18594855	227	543	642	1350148
3	1	4260473	34048	175983	155050	1350148
4	2	2977	1274	2570	26091	1350148

3.0.2 Avazu

The **Avazu** dataset was originally released in 2014 for a CTR prediction Competition on Kaggle (Cite Avazu). The data is composed of 11 days worth mobile ad marketplace data. Much like the KDD12 dataset above, this dataset contains features ranging from user activity (clicks), user identification (device type, IP) to ad features. Notable differences to the KDD12 dataset include the fact that Avazu contains an “hour” feature (enabling the establishment of sequentiality of behaviours) and the fact that Avazu does not seem to contain query and ad texts.

```
[3]: %run -i scripts/data_analysis_and_preprocessing/retrieve_avazu.py
```

Snapshot of Avazu training data:

	id	click	hour	c1	banner_pos	site_id	\
0	15674134821169810910	1	14102300	1005	0	85f751fd	
1	15674278914362889244	0	14102300	1005	0	85f751fd	
2	1567455966106046075	0	14102300	1005	0	26fa1946	
3	15674616734887926359	0	14102300	1005	0	85f751fd	
4	15674670592044781339	0	14102300	1005	0	85f751fd	

	site_domain	site_category	app_id	app_domain	...	device_type	\
0	c4e18dd6	50e219e0	e71aba61	2347f47a	...	1	
1	c4e18dd6	50e219e0	6f8bcb0f	2347f47a	...	1	
2	e2a5dc06	3e814130	ecad2386	7801e8d9	...	1	
3	c4e18dd6	50e219e0	53de0284	d9b5648e	...	1	
4	c4e18dd6	50e219e0	a0fc55e5	2347f47a	...	1	

	device_conn_type	c14	c15	c16	c17	c18	c19	c20	c21
0	0	21676	320	50	2495	2	167	-1	23
1	0	20476	320	50	2348	3	427	100005	61
2	0	20362	320	50	2333	0	39	-1	157
3	0	21611	320	50	2480	3	297	100111	61
4	0	20361	300	250	2333	0	39	-1	157

[5 rows x 24 columns]

3.0.3 Criteo

Finally, the Criteo dataset is another benchmark CTR prediction dataset that was originally released on Kaggle for a CTR prediction competition. The original dataset is made up of 45 Million user’s click activity, and contains the click/no-click target along with 26 categorical feature fields and 13 numerical feature fields. Unlike the other two datasets however, the semantic significance of these fields is not given - they are simply labelled as “Categorical 1-26” and “Numerical 1-13” respectively.

```
[4]: %run -i scripts/data_analysis_and_preprocessing/retrieve_criteo.py
```

Snapshot of Criteo training data:

	click	int_1	int_2	int_3	int_4	int_5	int_6	int_7	int_8	int_9	\
0	0	NaN	1	2.0	5.0	27586.0	32.0	2.0	14.0	21.0	

1	1	14.0	1	1.0	8.0	276.0	14.0	41.0	9.0	10.0
2	0	NaN	1	27.0	25.0	NaN	NaN	0.0	54.0	55.0
3	0	0.0	442	1.0	1.0	3029.0	58.0	2.0	13.0	44.0
4	0	0.0	-1	2.0	1.0	1167.0	88.0	23.0	19.0	673.0

	...	cat_17	cat_18	cat_19	cat_20	cat_21	cat_22	cat_23	\
0	...	07c540c4	bdc06043	NaN	NaN	6dfd157c	NaN	32c7478e	
1	...	e5ba7672	87c6f83c	NaN	NaN	0429f84b	NaN	be7c41b4	
2	...	2005abd1	87c6f83c	NaN	NaN	15fce809	NaN	be7c41b4	
3	...	d4bb7bd8	cdfa8259	NaN	NaN	20062612	NaN	dbb486d7	
4	...	27c07bd6	5bb2ec8e	49b8041f	b1252a9d	bff87997	NaN	32c7478e	


		cat_24	cat_25	cat_26
0		ef089725	NaN	NaN
1		c0d61a5c	NaN	NaN
2		f96a556f	NaN	NaN
3		1b256e61	NaN	NaN
4		3fdb382b	f0f449dd	49d68486

[5 rows x 40 columns]

3.0.4 Target Variable Analysis

The figure below shows that the three datasets have vastly different average Click Through Rates per instance. The average CTR for the KDD12 dataset is only 3.4%, whereas the Criteo dataset is 25.6%.


```
[5]: %run -i scripts/data_analysis_and_preprocessing/ctr_bar_charts.py
```



3.0.5 Missingness and Data Imputation

Below, I take a look at whether or not our dataset has any missing values.

```
[6]: # Show missingness matrix for criteo dataset
Image(filename='figures/criteo_missingno_matrix.png')
```



Above we see that Criteo has some missing values. Below I proceed by imputing the missing values using Sklearn's KNN Imputer. The code for these imputations does not get executed here.

Imputation steps taken were:

1. Factorize categorical values in the dataset, converting them to integers. This was done because sklearn's imputers only work with numerical data.
2. Use the sklearn's `IterativeImputer` with `HistGradientBoostingRegressor` to impute the missing values. This was recommended in [this github discussion](#) for imputing data with categorical values.
3. Concatenate missingness indicators to the dataset as additional features, as recommended by Van Buuren (2018)


The script for the above is in `scripts/data_analysis_and_processing/impute_criteo_nulls.py`.

```
[7]: # Pick up the imputed dataset
criteo_imputed_inds = pd.read_csv('./data/criteo/criteo_train_imputed.csv').
    ↪ astype('int')
criteo_imputed_inds[criteo.columns[criteo.dtypes == 'category']] =_
    ↪ criteo_imputed_inds[criteo.columns[criteo.dtypes == 'category']].
    ↪ clip(lower=0, upper=None)
```

3.0.6 Sparse Multi-Value Categorical Features

As already mentioned above, ad marketplace data often contains sparse categorical features, which make signal detection extremely difficult in shallow modelling frameworks. Below I show examples from each dataset

```
[8]: %run -i scripts/data_analysis_and_preprocessing/plot_cardinalities.py
```



A common remedy to the above issue is to *bin* the categorical feature values before one-hot encoding or embedding, according to some given threshold (Cite Song, Others). This essentially means that for a given threshold t , we retain only the values for the multi-value categorical features that have more than t occurrences in the dataset. (Cite Song) Recommends usign, setting $t = 10, 5, 10$ for Criteo, KDD12 and Avazu respectively. Due to computational limitations, this was multiplied by a factor of 100

```
[10]: %run -i scripts/data_analysis_and_preprocessing/binning_OH_encoding.py
```

Before one-hot encoding:

KDD12 shape: (100000, 12)

Avazu shape: (100000, 24)

Criteo shape: (100000, 64)

After one-hot encoding:

KDD12 shape: (100000, 12)

Avazu shape: (100000, 24)

Criteo shape: (100000, 64)

Sparse output:

KDD12 shape: (100000, 1478)


Avazu shape: (100000, 924)

Criteo shape: (100000, 2003)

3.0.7 High Variance Numerical outliers


Below I check the distributions of the numerical features in the datasets

```
[11]: %run -i scripts/data_analysis_and_preprocessing/plot_numerical_distributions.py
```



Due to the high variance of numerical features in the Criteo dataset, it is necessary to transform these variable in order to ease the training of deep NN's. As done by (Cite Song and Wang, and the winner of the Criteo Competition), we will proceed by applying the transform $\log^2(z)$ if $z > 2$, and where z is the standardized numerical value.

```
[12]: %run -i scripts/data_analysis_and_preprocessing/numerical_standardization.py
```




```
[13]: %run -i scripts/data_analysis_and_preprocessing/export_preprocessed.py
```

3.0.8 Correlation Analysis

In the [following script](#), conduct a correlation analysis of the features to the Click-Through rate.

```
[14]: Image(filename = './figures/criteo_corr_matrix.png')
```

```
[14]:
```

Some very high correlations between some of the features across fields in all three datasets. This possibly points to there being potential for dimensionality reduction across this feature set.

Unfortunately, from the correlation heatmaps, there appears to be little to no correlation between the first-order features and the target click variable.

4 Modelling

In this section I will compare the performance of two shallow modelling approaches (Logistic Regression and Factorization Machines) to a naive DNN for CTR prediction. As with (Cite Song and Wang), I will use the **Area Under the ROC Curve** and **Logloss** measures to compare the performance of the different modelling approaches on the test set.

```
[15]: %run -i scripts/modelling/load_and_prep_data.py
```

4.1 Logistic Regression

```
[16]: %run -i scripts/modelling/fit_lr_models.py
```

```
[17]: %run -i scripts/modelling/score_lr_models.py
```

KDD12:

Log loss: 0.1624932293410118

ROC AUC: 0.69918173566772

Accuracy: 0.95825

Avazu:

Log loss: 0.41220025365942664

ROC AUC: 0.7187741663639018

Accuracy: 0.83205

Criteo:

Log loss: 0.4934802502304099

ROC AUC: 0.7450126464008422

Accuracy: 0.7671

```
[18]: %run -i scripts/modelling/save_lr_models.py
```

4.2 Factorization Machine

Below I proceed by applying the SGD solver, as shown in the [relevant tutorial](#) for FastFM (Cite).

Note: Unfortunately, the fastFM library is currently only compatible with Linux and iOS. Since I have a Windows PC, I ran the training script below on an AWS Sagemaker Instance.

```
[19]: %run -i scripts/modelling/fit_fm_models.py
```

```
[20]: %run -i scripts/modelling/score_fm_models.py
```

KDD12:

Log loss: 0.31685987446578695

ROC AUC: 0.5336269490760196

Accuracy: 0.95825

Avazu:

Log loss: 10.974548402513749

ROC AUC: 0.5212335152327499

Accuracy: 0.66185

Criteo:

Log loss: 15.774504905747122

ROC AUC: 0.47794380839584155
Accuracy: 0.56235

```
[21]: %run -i scripts/modelling/save_fm_models.py
```

4.3 MLP

In this section, I implement a simple 3 layer MLP for CTR prediction.

```
[22]: %run -i scripts/modelling/load_tf_datasets.py
```

```
2024-06-16 20:55:29.331209: W
external/local_tsl/tsl/framework/cpu_allocator_impl.cc:83] Allocation of
1281280000 exceeds 10% of free system memory.
```

```
[23]: %run -i scripts/modelling/fit_kdd12_mlp.py
%run -i scripts/modelling/fit_avazu_mlp.py
%run -i scripts/modelling/fit_criteo_mlp.py
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	189056
batch_normalization (Batch Normalization)	(None, 128)	512
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 64)	8256
batch_normalization_1 (Batch Normalization)	(None, 64)	256
dropout_1 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 32)	2080
batch_normalization_2 (Batch Normalization)	(None, 32)	128
dropout_2 (Dropout)	(None, 32)	0
dense_3 (Dense)	(None, 1)	33

Total params: 200321 (782.50 KB)
Trainable params: 199873 (780.75 KB)


Non-trainable params: 448 (1.75 KB)

None

2024-06-16 21:00:24.977175: W
external/local_tsl/tsl/framework/cpu_allocator_impl.cc:83] Allocation of
1281280000 exceeds 10% of free system memory.

[]:

[25]: %run -i scripts/modelling/plot_mlp_training.py



[27]: %run -i scripts/modelling/eval_mlp_models.py

KDD12:
200/200 [=====] - 1s 2ms/step - loss: 0.1769 -
accuracy: 0.9500 - auc: 0.6809

Avazu:
200/200 [=====] - 0s 2ms/step - loss: 0.4156 -
accuracy: 0.8317 - auc: 0.7192

Criteo:
200/200 [=====] - 1s 2ms/step - loss: 0.5081 -
accuracy: 0.7657 - auc: 0.7428

5 Summary of findings

6 Suggested Future Research

7 References

- eMarketer. (2023). Digital advertising spending worldwide from 2021 to 2027 (in billion U.S. dollars) . Statista. Statista Inc.. Accessed: June 09, 2024. <https://www-statista-com.iclibezp1.cc.ic.ac.uk/statistics/237974/online-advertising-spending-worldwide/>
- Aden, Yi Wang. (2012). KDD Cup 2012, Track 2. Kaggle. <https://kaggle.com/competitions/kddcup2012-track2>
- Steve Wang, Will Cukierski. (2014). Click-Through Rate Prediction. Kaggle. <https://kaggle.com/competitions/avazu-ctr-prediction>
- Jean-Baptiste Tien, joycenv, Olivier Chapelle. (2014). Display Advertising Challenge. Kaggle. <https://kaggle.com/competitions/criteo-display-ad-challenge>
- Van Buuren, S. (2018). Flexible imputation of missing data. CRC press.
- Song, W., Shi, C., Xiao, Z., Duan, Z., Xu, Y., Zhang, M., & Tang, J. (2019, November). Autoint: Automatic feature interaction learning via self-attentive neural networks. In Proceedings of the 28th ACM international conference on information and knowledge management (pp. 1161-1170).
- Wang, F., Gu, H., Li, D., Lu, T., Zhang, P., & Gu, N. (2023, October). Towards Deeper, Lighter and Interpretable Cross Network for CTR Prediction. In Proceedings of the 32nd ACM International Conference on Information and Knowledge Management (pp. 2523-2533).